



COMS3010A

Operating Systems

Project 1 - My Alloc

Instructors

COMS3010A Lecturer:

Branden Ingram branden.ingram@wits.ac.za

COMS3010A Tutors:

Guy Axelrod	1864180
Molefe Molefe	1858893
Mihai Branga-Peicu	1833986
Tshegofatso Mohlala	1940009
Bongumusa Mavuso	1682836
Uwais Suliman	1825613
Jonas Chirindza	1900596
Kaone Senoelo	1819369
Khumo Tsagae	1929633
Tarshen Naidoo	1659311
Jared Harris-Dewey	1846346
Oluwademilade Osikoya	1963195
Nhlanhla Mpele	1860226
Thabo Ranamane	1831661
Sakhile Mabunda	1830088

Consultation Times

Questions should be firstly posted on the moodle question forum, for the Lecturer and the Tutors to answer. If further explanation is required consultation times can be organised.

1 Introduction

This is an assignment where you will implement your own malloc using the buddy algorithm to manage the splitting and coalescing of free blocks. You should be familiar with the role of the allocator and the theory behind the operation of the buddy algorithm from the Lectures. In this project you will not be given all details on how to do the implementation, but the general strategy will be outlined. You will be provided with some helper functions which you will be required to use in order to implement the buddy allocator. A Project Template has been provided on moodle to use. The functioning of these helper functions is explained below.

2 The buddy algorithm

The buddy memory allocation technique is a memory allocation algorithm that divides memory into partitions to try to satisfy a memory request as suitably as possible. This system makes use of splitting memory into halves to try to give a best fit. The benefit of the buddy algorithm is that the amortized cost of the allocate and free operations are constant. The slight down side is that we will have some internal fragmentation since blocks

only come in powers of two. We will keep the implementation as simple as possible, and therefore we will not do the uttermost to save space or reduce the number of operations. The idea is that you should have a working implementation and do some benchmarking which you can later optimise upon.

2.1 Implementation Guidelines

Your implementation should follow the guidelines outlined in this document. In this implementation the largest block that we will be able to handle is 4KB. We can then split it into sub-blocks and the smallest block will be 32 bytes (2^5). This means that we will have 8 different sizes of blocks: 32bytes, 64, 128 , .. 4 KB.

We will refer to these as a level-0 block, a level-1 block etc. We encode these as constants in our system:

- MIN the smallest block is (2^{MIN})
- LEVELS the number of different levels
- PAGE the size of the largest block ($2^{12} = 4KB$)

When we allocate a new block of 4KB, it needs to be aligned on a 4KB boundary i.e. the lower 12 bits are zero. This is important and we will get this for free if we choose the system page sizes when allocating a new block

2.2 operations on a block

A block consists of a head and a byte sequence. It is the byte sequence that we will hand out to the requester. We need to be able to determine the size of a block, if it is taken or free and, if it is free, the links to the next and previous blocks in the list. You will find a “struct” defined to represent the head. The size of the block is represented by a variable called “level” and is encoded as 0; 1; 2 etc. This is the index in the set of free lists that we will keep. Index 0 is for the smallest size blocks i.e. 32 byte. a variable called “status” to indicate if the memory region is free or allocated as well as variables called “next” and “prev” for our block pointers.

```
struct head {
    enum AEflag status;
    short int level;
    struct head *next;
    struct head *prev;
}

enum AEflag{Free, Taken};
```

A set of functions are required to implement the functionality of the buddy algorithm. Some of these have been provided in the project template. Given these functions you should be able to implement the algorithm independent of how we choose to represent the blocks. **Ultimately you will be required to implement the following functions:**

- new - Section [2.2.1](#)
- level - Section [2.2.2](#)
- bfree - Section [2.4.2](#)
- balloc - Section [2.4.1](#)

while making use of the following functions:

- buddy
- split
- primary
- magic
- hide

2.2.1 a new block

To begin with you have to create new block. you can do this using the `mmap()` system call. This procedure will allocate new memory for the process and here we allocate a full page i.e. a 4KB segment. Here you can take for granted that the segment returned by the `mmap()` function is in fact aligned on a 4 KB boundary. Make sure to initialise the status and level variables as well include a failure check for `mmap` using `MAP_FAILED`.

```
struct head *new() {
// You are required to complete
//
//
//
}
```

Look-up the man pages for `mmap()` and see what the arguments mean. If you extend the implementation to handle larger blocks you will have to change these parameters.



Figure 1: `struct head* block = new();`

This function should as demonstrated above return a pointer to the start of a `struct head*` object

2.2.2 level

If the user request a certain amount of bytes we need a slightly larger block to hide the head structure. When we have found a size that is large enough to hold the total number of bytes we know the level.

```
int level(int req) {
// You are required to complete
//
//
//
}
```

2.3 Helper Functions

The following subsections are explanations to the functioning of the helping functions which have been proved to you. You must not change these in any way.

2.3.1 who's your buddy

Given a block we want to find the buddy. The buddy of the block is the other pair of that block.




Figure 2: `struct head* tmp's buddy = buddy(tmp);`



Figure 3: `struct head* tmp's buddy = buddy(tmp);`




Figure 4: `struct head* tmp's buddy = buddy(tmp);`

The figures above demonstrate the desired behaviour of this function. One way of doing this is by toggling the bit that differentiate the address of the block from its buddy. For the 32 byte blocks, that are on level 0, this means that we should toggle the sixth bit. If we shift a 1 five positions (MIN) to the left we have created a mask that we can xor against the pointer to the block.

```
struct head *buddy(struct head* block) {  
    int index = block->level;  
    long int mask = 0x1 << (index + MIN);  
    return (struct head*) ((long int) block ^ mask);  
}
```

2.3.2 split a block

When we don't have a block of the right size we need to divide a larger block in two.




Figure 5: `struct head* tmp's split = split(tmp);`

This can be achieved by first finding the level of the block, subtract one and create a mask that is or'ed with the original address. This will now give us a pointer to the second part of the block.

```

struct head *split(struct head* block) {
    int index = block->level-1;
    int mask = 0x1 << (index + MIN);
    return (struct head*) ((long int)block | mask);
}

```

2.3.3 the primary of two blocks

We also need a function that merges two buddies. When we are to perform a merge, we first need to determine which block is the primary block i.e. the first block in a pair of buddies. The primary block is the block that should be the head of the merge block so it should have all the lower bits set to zero.

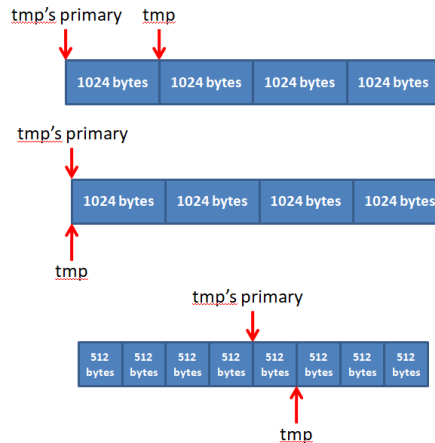


Figure 6: `struct head* tmp's primary = primary(tmp);`

We achieve this by masking away the lower bits up to and including the bit that differentiate the block from its buddy. Note that it does not matter which buddy we choose, the function will always return the pointer to the primary block.

```

struct head *primary(struct head* block) {
    int index = block->level;
    long int mask = 0xffffffffffffffff << (1 + index + MIN);
    return (struct head*) ((long int)block & mask);
}

```

2.3.4 the magic

We use the regular magic trick to hide the secret head when we return a pointer to the application. We hide the head by jumping forward one position. If the block is in total 128 bytes and the head structure is 24 bytes we will return a pointer to the 25'th byte, leaving room for 104 bytes.




Figure 7: `struct head* tmp's hide = hide(tmp);`

```

struct head *hide(struct head* block) {
    return (void*) (block +1);
}

```

The trick is performed when we convert a memory reference to a head structure, by jumping back the same distance

```

struct head *magic(void *memory) {
    return ((struct head*)memory-1);
}

```

2.3.5 before you continue

Implement the above functions in a file called buddy.c. In a file test.c there is a main() procedure that calls the test procedure. You will find some basic testing functions however, these are nowhere near sufficient and only serve as an example. Do extensive testing, if anything is wrong in these primitive operations it will be a nightmare to debug later. Also implement a function that runs your own tests that ensures you that the operations work as expected. You can for example create a new block, divide it in two, divide in two and then hide its header, find the header using the magic function.

2.4 the algorithm

So given that the primitive operations work as intended, it is time to implement the algorithm. Your task is now to implement two procedures balloc() and bfree() that will be the API of the memory allocator. We do not replace the regular malloc() and free() since we want to use them when we write our benchmark program. We will use one global structure that holds the double linked free lists of each layer.

```

struct head flists [LEVELS] = {NULL} ;

```

Our flists is basically an array of linked lists where each track the blocks of a specific size. Here since we have only initialised our new block we have one entry in the list at index 7.

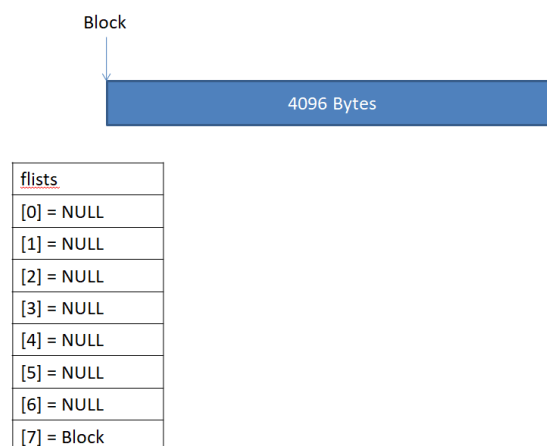


Figure 8: Flists Example

2.4.1 balloc

The first part of our buddy memory allocator will be the allocation component. Below is the function structure as well as some psuedocode

```
void *balloc(size_t size) {
//Check if size is 0
//Determine required level based upon size
//Based on that level do the necessary splitting
//  if required remembering to update flists and
//  return a pointer to the allocated memory block
//Finally hide the header and return the pointer
}
```

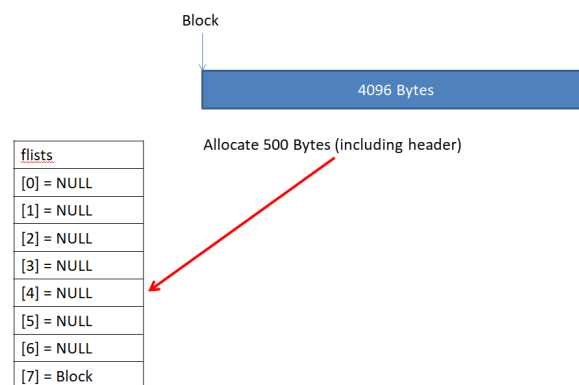


Figure 9: Allocation Example

If we were to allocate 500 Bytes we would need to first determine the index corresponding to the smallest block required to store this amount of data. In this case it would be 4. Since flists[4] is Null we need to split in order to divide our memory into the required block size. If there was already a free block in this list we would only need to return a pointer to it.

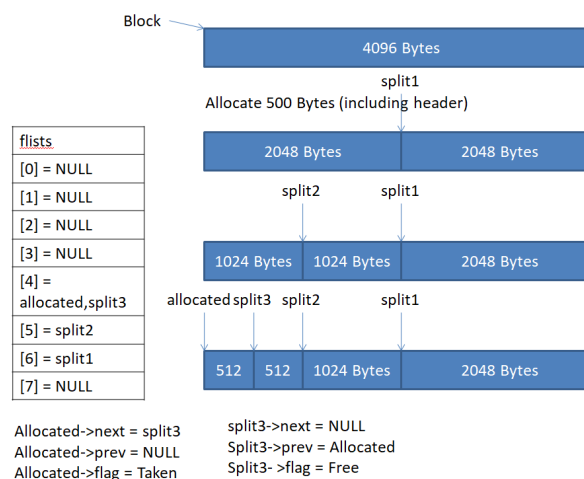


Figure 10: Allocation Example

Here you can see some information demonstrating the process of splitting that would be required to find a pointer to a block of the right size. Additionally you can see the state of flists. Note how it no longer contains an item in flists[7] since that block is no longer available. It is up to you to decide in which way you add or remove from the linked lists.

2.4.2 bfree

The second part of our buddy memory allocator will be the free component. Below is the function structure as well as some pseudocode

```
void bfree(void *memory) {
//Check memory is a null pointer
//Get the header of the pointer
//Update flists and free the block which the pointer points to
//Perform merging up the list if possible
}
```




Figure 11: Free Example

To free some memory we need to identify which part of flists it belongs to.




Figure 12: Free Example

We then have to update that block in terms of its status.

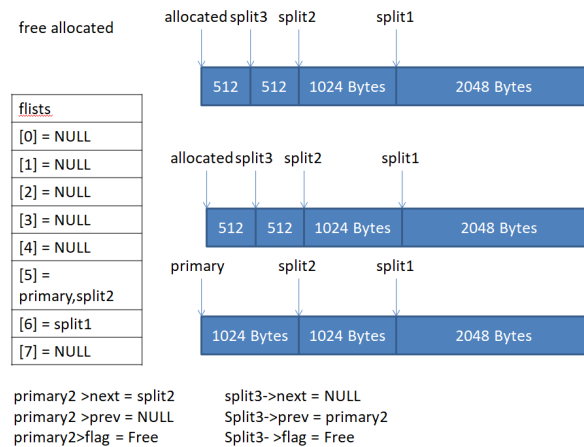


Figure 13: Free Example

lastly if merging is possible, i.e. if its buddy is also free, we need to combine these two blocks and update flists accordingly. If this results in the possibility to merge further up the tree you need to do so.

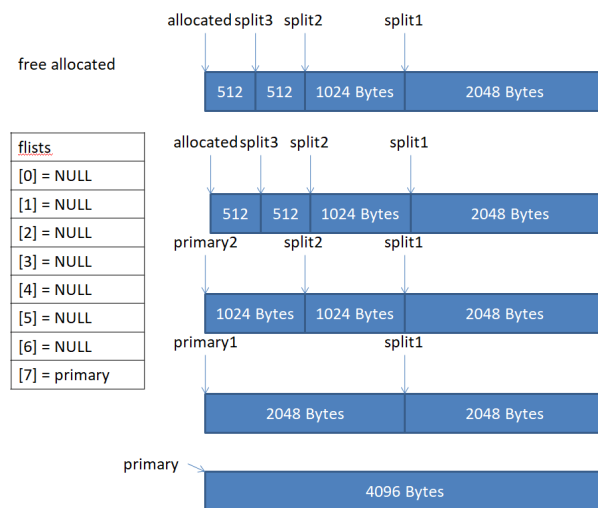


Figure 14: Free Example

Draw pictures that describes how the double linked list is manipulated. Do small test as you proceed, make sure that simple things work before you continue.

3 Submission

I am trying to set up an automatic moodle marker, however, at the moment it does not support C. In the meantime you should run extensive testing locally to determining if your functions have been implemented correctly. Lastly there will be a small report component in the form of a moodle quiz.

Academic Integrity

There is a zero-tolerance policy regarding plagiarism in the School. Refer to the General Undergraduate Course Outline for Computer Science for more information. Failure to adhere to this policy will have severe repercussions.

During assessments:

- You may not use any materials that aren't explicitly allowed, including the Internet and your own/other people's source code.
- You may not access anyone else's Sakai, Moodle or MSL account.

Offenders will receive 0 for that component, may receive FCM for the course, and/or may be taken to the legal office.